

Call Coach Lab v1 — Role-Based User Journey Bible

Audience: Engineers, designers, PM, interns

Purpose: Make it impossible to “miss” a flow, state, screen, or edge case while building v1.

Scope: v1 (AI Call Auditor + minimal “action loop” guidance).

Explicitly not in v1: mobile apps, multi-location workspaces, missed-call tracking, full AI learning plans, long-form coaching programs.

The three people this product must make happy

- **Admin (Owner/Ops Admin):** sets up the system, integrations, permissions, scorecards, and visibility. Wants ROI clarity + control.
 - **Manager (Sales/Front Desk Manager):** runs daily QA, ensures process adherence, improves the team, sees trends.
 - **Agent (Front desk / sales rep):** receives fair feedback with evidence, improves quickly, sets outcomes, disputes mistakes.
-

Part 1 — The world of CallAudit (things everyone must agree on)

1) The “units” of the product (Core Objects)

If the team agrees on these, almost everything else becomes straightforward.

Workspace

- One business (one clinic/branch) = one workspace.
- No multi-location in v1.

User

- Belongs to one workspace.
- Has one role: **Admin, Manager, Agent.**
- Can optionally belong to a **Team** (Front Desk, Sales, Follow-up, etc.).

Team

- A grouping mechanism for agents.
- Used for filtering, reporting, and assignment.
- Optional, but recommended.

Call

The fundamental artifact we score and review. Every call has:

- Audio recording (required in v1)
- Transcript (generated)
- Metadata (some required, some optional)
- A scorecard attached
- A score + per-criterion results
- Flags (AI-detected)
- A call outcome (set by agent; can be flagged by AI)

Contact

This is the **customer journey object**.

- **Primary key in v1:** normalized phone number (E.164 style).
- A contact contains multiple calls over time (inbound lead + follow-ups + consult calls).
- Contact also contains the **Final Outcome** (journey outcome), which is not tied to a single scorecard.

Why Contact exists: you explicitly need to see the complete journey across calls (same phone number) and tie it to the final outcome. Without Contact, ROI and conversion reporting becomes misleading.

Scorecard

A structured set of criteria used to audit calls.

- Multiple scorecards per workspace (Inbound / Follow-up / Consult etc.)
- Scorecards have:
 - sections
 - criteria
 - scoring rules (points/weights)
 - pass/fail thresholds
 - “critical fails”
 - versioning (Draft vs Active)
 - outcomes set (per scorecard)

Outcome Sets

There are **two outcome concepts** in v1:

A) Call Outcome (per scorecard)

- Set by the **agent** per call.
- Fully custom per scorecard.
- Example (Inbound scorecard):
 - Booked Consult
 - Follow-up Scheduled
 - Not Interested
 - Not Fit
 - Spam

B) Final Outcome (per contact journey)

- A single “final truth” about the contact journey.
- Also fully custom per business (because clinics define conversion differently).
- Example final outcomes:
 - Booked
 - Converted / Paid
 - Lost
 - Not Fit
 - Spam

How do these relate?

We need a simple mapping layer: **Call Outcome** → **Journey Stage/Final Outcome impact** (configurable).

- Example:
 - Inbound “Booked consult” → sets contact stage to “Booked”
 - Consult “Paid” → sets contact final outcome to “Converted”
 - Follow-up “No response” → does not set final outcome, but keeps stage “In progress”

This mapping prevents flawed reporting when outcomes are per scorecard.

Flags (AI findings)

Flags are “signals” the system detects from transcript/audio:

- Objection flags: Price, competitor, insurance, time, spouse approval
- Compliance flags: refund promise, medical claims, sensitive info
- Outcome mismatch: “Agent marked converted but customer refused”
- Spam risk / wrong number detection

- Escalation triggers: lawsuit, chargeback, etc.

Evidence

Every important claim in the product must point to:

- transcript snippet + timestamp
 - playable audio at that moment
-

2) Non-negotiables for v1 (so the product isn't "AI vibes")

These are the rules that prevent trust collapse:

1. **Every score must be explainable**
Pass/Fail/Unclear + confidence + timestamp evidence.
 2. **Unclear is allowed**
If evidence is missing, the model must not guess.
 3. **Scoring SLA: < 5 minutes**
From call ingest → transcript + score available.
This impacts architecture + UX states.
 4. **Agent sets call outcome, AI can flag it**
Manager cannot realistically set outcomes for every call.
 5. **Call type can be AI guessed, manually overridden**
AI guess uses contact history + call content.
 6. **Works in Hindi + English + Punjabi**
At minimum: transcription + keyword search + scoring must handle these.
-

3) The end-to-end loop (the thing investors and interns should remember)

Calls come in → get scored → reviewed if needed → agent improves → trends improve → ROI becomes visible.

In product terms:

1. Ingest call (Upload / Exotel / Twilio)
2. Create transcript + speaker labels
3. AI assigns call type + scorecard (guess)
4. AI scores with evidence + confidence

5. Call is routed:
 - Auto-pass (high confidence, no critical fail)
 - QA Queue (unclear / low score / triggers / new rep / random sample)
 6. Manager reviews & overrides (if needed)
 7. Agent receives evaluation + acknowledges
 8. Outcome is set (agent)
 9. Contact timeline updates (journey)
 10. Insights update (by source, by rep, by scorecard, by outcome)
-

Part 2 — Admin Journey (Owner/Ops Admin)

Admin storyline: “Aman sets up SmileBright Dental”

Aman doesn't want software. Aman wants **less waste**, **more bookings**, and **less dependence on star reps**.

Your v1 experience must get Aman to “first scored call” fast, without asking 50 questions.

Admin Journey Phase 0 — Getting into the workspace

Entry paths

Path A: Admin signs up

- Admin creates workspace
- Enters business name + business type (Dental / Skin / Hair / Chiro / Other)
- Chooses default language preference (multi-language enabled by default)

Path B: Admin invited

- Admin clicks invite link
- Sets password
- lands in setup checklist

Required design states

- Invite expired

- Invite already used
 - Wrong email
 - Password rules + error messaging
-

Admin Journey Phase 1 — Workspace setup wizard

Wizard steps (v1)

1. Business basics
2. Teams & users
3. Connect calls
4. Scorecards + outcomes
5. Launch baseline

Step 1: Business basics

Admin inputs:

- Workspace name
- Industry type
- Timezone (Asia/Kolkata default but editable)
- Default language preferences (Hindi/English/Punjabi enabled)

UX detail:

The wizard must explain *why* we ask each thing in one line:

- “Timezone ensures call timestamps match your day.”
 - “Language helps transcript accuracy.”
-

Admin Journey Phase 2 — Teams & users

Create teams

- “Front Desk”
 - “Consultants”
 - “Follow-up”
- (Teams are optional but recommended)

Invite users

Admin invites:

- Managers (1–N)
- Agents (1–N)

Invite contains:

- Role
- Team assignment (optional)
- Welcome text (auto)
- Link expiry (configurable)

Permissions control (your requirement):

Admin can toggle:

- ☒ Managers can edit scorecards
 - ☒ Managers can edit outcome sets
 - ☒ Managers can publish scorecards
 - ☒ Managers can manage integrations
- (Each toggle has a tooltip explaining impact)

UX details required:

- bulk invite (paste emails)
- resend invite
- revoke invite
- change role after invite accepted
- deactivate user (without deleting data)
- team reassignment

Admin Journey Phase 3 — Connect calls (Upload + Exotel + Twilio)

Option 1: Upload (Bulk) — Audio + CSV (recommended)

This is a primary v1 path.

Upload requirements

Admin/manager/agent uploads:

- One or many audio files
- One CSV file that includes per call row mapping

CSV schema (v1)

Must include (minimum viable):

- `file_name` (must match audio filename exactly)
- `agent_name` (or `agent_email`)
- `call_datetime` (timezone assumed or specified)
- `call_type` (optional if AI guesses, but recommended)
- `source` / `channel` (optional)
- `campaign` / `utm_campaign` (optional)
- `caller_number` (recommended for contact journey)
- `receiver_number` (optional)
- `language` (optional; else auto-detect)
- `call_outcome` (optional; agent can set later)

UX must include:

- Download CSV template button
- “Validate CSV” step before processing
- Error report with row-level issues:
 - missing `file_name`
 - file not found
 - unknown agent
 - invalid datetime
 - duplicate file name
- Partial success allowed:
 - Process valid rows
 - Reject invalid rows with clear reasons

Upload flow (screen-by-screen)

1. Upload page shows:
 - “Drop audio files here”
 - “Upload CSV mapping”
2. System validates:
 - file count vs CSV rows
 - matching filenames
3. System shows “Ready to process: X calls”
4. Click “Start processing”
5. Processing queue appears with:
 - transcription progress
 - scoring progress
 - failures list
6. Completion summary:

- processed
- failed
- needs review (unclear)

Edge cases to handle

- Audio too long (limit defined)
 - Corrupted audio
 - Multiple audio files same name
 - CSV includes calls without audio
 - Audio exists without CSV row
 - User uploads CSV first then audio
-

Option 2: Exotel integration (v1)

Admin connects Exotel to ingest recordings automatically.

Integration setup steps

1. Admin opens Settings → Integrations → Exotel
2. Click “Connect”
3. Enters credentials / API key (per Exotel docs)
4. Select ingest mode:
 - Pull recordings periodically OR webhook-driven (implementation choice)
5. Map fields:
 - Agent mapping strategy (depends on Exotel data):
 - if agent extension available → map extension to agent user
 - if agent name available → map by name
 - else → calls come in “Unassigned” and manager assigns

Mandatory fallback in v1:

If agent identity cannot be reliably detected, calls should still ingest and appear in **Calls Inbox** → “**Unassigned agent**” filter.

UX states required

- Connected (green)
 - Permission denied / invalid credentials
 - No recordings found yet
 - Recording fetch failed (retry)
 - Rate limit handling message (non-scary)
-

Option 3: Twilio integration (v1)

Same architecture as Exotel, different fields.

Setup steps

- Settings → Integrations → Twilio
- Connect with credentials
- Choose which phone numbers to ingest
- Recording types:
 - mixed/mono
 - dual-channel (if available)

Speaker labeling behavior in v1:

- If dual-channel: use it for speaker separation
 - If mono: diarization + “Agent vs Customer” inference
-

Admin Journey Phase 4 — Attribution setup (UTM/campaign + manual)

Your requirement: **connect call with UTM/campaign payload from wherever the client captures it** + manual option.

v1 approach (practical + flexible)

We implement a **Lead Attribution Resolver** with multiple inputs:

Attribution input methods (v1):

1. From Call Upload CSV (best)
2. From integration metadata (if Exotel/Twilio provides custom params)
3. From a “Lead Metadata Import” CSV (separate upload):
 - phone number
 - lead created timestamp
 - utm_source, utm_medium, utm_campaign, gclid, campaign name
4. Manual assignment per call (fallback)

Matching logic in v1 (explicit):

- If call already has source/campaign → keep it
- Else, try to match lead metadata:
 - same phone number

- lead created within configurable time window (e.g., 7 days before call)
- If still unknown → mark as “Unknown source” and allow manual set

UX requirements

- A “Sources” settings page:
 - define channels (Google Ads, Meta, Organic, Referral...)
 - optional campaign list
- A “Lead import” tool:
 - upload lead metadata CSV
 - view match rate (X% calls matched)
 - view unmatched calls

Important:

ROI in v1 is **trends/insights by source/channel**, not cost.

Admin Journey Phase 5 — Scorecards + outcome sets (multiple scorecards)

Admin configures:

- Scorecard templates and/or custom scorecards
- Multiple scorecards
- Outcomes per scorecard
- Contact final outcomes (journey outcomes)
- Mapping rules (call outcome → contact stage)

Scorecard setup flow (v1)

1. Scorecards → Templates
2. Choose template (Inbound Lead, Follow-up, Consult)
3. Import as **Draft**
4. Edit scorecard in Scorecard Builder
5. Define:
 - threshold
 - critical fails
 - criteria evidence requirements
6. Define **Call Outcomes** for this scorecard
7. Publish scorecard (Active)

Outcome setup flow (v1)

1. For each scorecard: create/edit call outcomes
2. Global: create/edit contact final outcomes
3. Configure mapping rules:
 - “If call outcome == X, set contact stage to Y”
 - “If call outcome == Z, set contact final outcome to Lost”

Why mapping matters:

It makes reporting correct across multi-call journeys.

Admin Journey Phase 6 — Launch baseline

Admin is done when:

- Calls are flowing in
- Scorecards are active
- First 10 calls scored
- Manager can open QA queue
- Agent can see evaluations

“Launch checklist” UI requirements

-  Integrations connected OR upload completed
 -  Users invited (at least 1 manager + 1 agent)
 -  At least 1 scorecard published
 -  At least 10 calls processed
 -  First QA review completed
-

Part 3 — Manager Journey (Daily operator + quality enforcer)

Manager storyline: “Priya runs QA without drowning”

Priya’s nightmare is listening to 50 calls/day.

Her dream is: **open one queue**, review in minutes, and improve bookings.

Manager Journey Phase 0 — First login

Manager accepts invite → lands on **Home Dashboard**.

If setup is incomplete

Home shows:

- “Setup incomplete” banner with action buttons:
 - connect calls (if not done)
 - publish scorecard
 - invite agents

If setup is complete

Home shows:

- Today snapshot:
 - Qualified calls
 - Booked consults
 - Avg QA score
 - Cost not needed in v1 (replace with “Qualified rate by source”)
 - “Today’s actions” panel:
 - X calls need review
 - Y unclear
 - Z critical fails
 - Trends:
 - Top missed criteria
 - Rep leaderboard
-

Manager Journey Phase 1 — Calls Inbox (triage + search)

Manager uses Calls Inbox to:

- find specific calls
- filter by team/agent/source/call type
- see score and flags
- open a call review

Calls Inbox must support (v1)

Filters:

- date range
- agent
- team
- call type
- scorecard
- source/campaign
- language
- flags
- “Unassigned agent”
- “Outcome mismatch flagged”

Columns (minimum):

- call date/time
- duration
- agent
- call type
- scorecard used
- score (0–100)
- outcome (agent-set) + mismatch badge
- top 1–2 misses
- flags
- action: open

States

- Processing (still scoring)
- Failed transcription (retry)
- Missing metadata (needs assignment)
- Unassigned agent (needs mapping)

Manager Journey Phase 2 — QA Queue (the daily workflow)

This is the page Priya opens every day.

Queue modes (v1)

- Lowest scores
- Unclear calls
- Critical fails
- New rep calls

- Random sample (model monitoring + fairness)
- Trigger terms (price, competitor, refund, lawsuit...)

Queue item summary (must show)

- Rep name + team
- call type
- scorecard
- score
- top 2 misses
- flags
- one-click “Jump to evidence moment” button

Manager daily routine (explicit flow)

1. Priya opens QA Queue
2. She starts with:
 - critical fails
 - unclear
 - lowest scores
3. Opens call review
4. Confirms whether scoring is correct
5. Overrides any wrong criterion
6. Adds short manager note
7. Sends a single next action to agent
8. Optionally saves a “best call clip” to library

Time target:

A manager should be able to review a call in **< 90 seconds** when everything is clear.

Manager Journey Phase 3 — Call Review (evidence-first)

Call review is where trust is built or destroyed.

Call Review layout (must-have)

- Audio player with playback speed
- Synced transcript
- Speaker labels (Agent vs Customer at minimum)
- Search in transcript (price, refund, competitor, etc.)
- Scorecard panel:
 - sectioned

- each criterion shows pass/fail/unclear + confidence + evidence snippet/time
 - override controls
- “Top misses” summary
- “One next action” box + send button
- Contact Journey tab/panel (see later)

Override logic (v1)

If Priya overrides:

- The call score recalculates
- The system stores:
 - original AI value
 - manager final value
 - override reason (optional but recommended)
- Overrides become training signals (future)

Trust features (v1)

- If AI is “Unclear,” show it prominently.
- If confidence is low, encourage review.

Manager Journey Phase 4 — Managing outcomes (without doing outcomes)

Agent sets call outcomes, but manager needs to ensure reporting is not polluted.

Outcome mismatch workflow (v1)

1. Agent sets “Converted”
2. AI detects transcript: “don’t call me”
3. System flags: **Outcome mismatch**
4. Manager sees this in:
 - QA queue filter “Outcome mismatch”
 - Calls inbox flag filter
5. Manager reviews quickly:
 - confirms mismatch
 - either:
 - requests agent correction (preferred)
 - or overrides outcome (if admin allows managers to override)

Outcome correction UX (v1)

- One-click “Request outcome correction”
 - Pre-filled note:
 - “Call indicates customer refused. Please update outcome.”
-

Manager Journey Phase 5 — Insights (weekly operator mode)

Manager uses Insights weekly, not hourly.

Insights tabs

1. Marketing ROI (trends by source/channel)
2. QA Trends
3. Voice of Customer

Marketing ROI (v1, no cost)

Show:

- call volume by source
- qualified rate by source (based on outcome + scoring gates)
- booked rate by source
- spam/not-fit rate by source
- top campaigns (if provided)
- “unknown source” share (should decrease as data improves)

QA trends (v1)

Show:

- avg score over time
- pass rate
- critical fails count
- top missed criteria
- rep variance
- section adherence (“golden scorecard” consistency)

Voice of Customer (v1)

Show:

- top objections and trends
 - lost reasons and trends
 - competitor mentions
 - pricing confusion themes
 - “suggested playbook update” (AI suggestion but should be editable by manager/admin)
-

Manager Journey Phase 6 — Team page (performance + actions)

Manager uses Team page to:

- spot who needs help
- nudge acknowledgements
- assign focus actions

Team page must include:

- agent list with:
 - avg score
 - pass rate
 - top miss
 - pending acknowledgements
 - last activity
 - right panel:
 - most missed criteria (team)
 - “assign next action to 3 reps”
 - pending acknowledgement nudges
 - new rep auto queue status
-

Part 4 — Agent Journey (the person who must actually change behavior)

Agent storyline: “Sarah wants to know what to do, not get judged”

Agents will resist if:

- scoring feels random
- feedback is vague
- they can't defend themselves

Agents will love it if:

- it's specific, fair, and fast
 - it helps them book more consults
 - it doesn't waste time
-

Agent Journey Phase 0 — First login

Agent accepts invite and lands on Agent Home.

Agent Home shows:

- score trend (7d)
- pass rate
- pending actions
- pending acknowledgements
- today's focus (most repeated miss)

Required agent setup elements

- profile name
 - assigned team
 - language preference (optional)
 - notification preferences (email + in-app)
-

Agent Journey Phase 1 — Calls happen

Calls come in via:

- Exotel/Twilio integrations (automatic)
- Upload by agent (manual)

Agent upload flow (single call)

Since you said agents can upload too, we must support non-CSV, non-bulk:

Agent clicks “Upload call” → sees:

- audio upload
- required fields:
 - call date/time (default: now)
 - call type (AI guess optional; can select)
 - scorecard (auto-assigned based on call type)
 - caller number (if known; else “unknown”)
 - source (optional)
- submit
- processing state appears in My Calls

Bulk upload remains CSV-first (manager/admin use case).

Agent Journey Phase 2 — Evaluation arrives (<5 minutes)

When scoring completes:

- Agent receives:
 - in-app notification: “New evaluation ready”
 - optional email: “Your call has been evaluated”
- My Evaluations shows new row at top

Evaluation list row must show

- call date/time + duration
 - score
 - pass/fail state (if threshold)
 - top misses (1–2)
 - next action summary (one line)
 - acknowledgement status (pending/acknowledged)
-

Agent Journey Phase 3 — Evaluation Detail (evidence-first)

Agent opens evaluation detail to see:

- score breakdown
- evidence clips

- manager note (if present)
- next action

What agent must be able to do (v1)

- Play audio at evidence timestamps
- Read transcript
- See why a criterion failed
- Acknowledge evaluation
- Request clarification (dispute) if wrong
- Set call outcome (if not already set)

Clarification workflow (v1)

Agent clicks “Request clarification”:

- chooses reason:
 - evidence missing
 - transcript incorrect
 - scoring incorrect
 - context issue
- adds comment
- submits

Manager receives a queue item: “Clarification requested”

Manager can:

- accept correction and override score
- reject with note
- mark “needs more info”

Agent Journey Phase 4 — Setting call outcome

Agent sets call outcome because manager can’t do it for every call.

Outcome setting UX (v1)

Where outcome can be set:

- end of evaluation detail (“Set outcome”)
- My Calls row quick action

Outcome field:

- uses outcome set from that call's scorecard

Outcome mismatch behavior (v1)

If agent sets outcome:

- AI checks if outcome aligns with transcript
- If mismatch, system:
 - shows warning to agent: "This outcome might not match what the customer said."
 - still saves (agent can override)
 - flags for manager review

This prevents "garbage outcomes" breaking reporting.

Agent Journey Phase 5 — The minimal action loop (v1)

This is where your product starts behaving like a coach, without becoming a full coach yet.

AI "next action" principles (v1)

- Must be **one thing** (not a list)
- Must be specific and derived from:
 - top misses
 - best practice ordering (pain → impact → price)
 - approved playbook guidance (if available)

Agent sees:

- "Next time, do this..."
- Example sentence
- Why it matters (one line)

No learning plan tracking in v1.

Just: "Here's the fix for your next call."

Part 5 — Contact Journey (your critical requirement)

Why we need a contact journey view

Because:

- multiple calls lead to conversion
 - outcomes per scorecard alone make reporting wrong
 - managers need to see the full timeline quickly
-

Contact Journey page (v1)

Where it lives:

- accessible from Call Review (manager) and Evaluation Detail (agent, limited)
- accessible from Calls Inbox by clicking the caller number/contact

Contact Journey must show

Header

- Phone number (primary identifier)
- optional name (if imported from leads)
- source/campaign (if known)
- current stage + final outcome
- last call date

Timeline

- each call card shows:
 - date/time
 - call type
 - scorecard used
 - score
 - call outcome (agent-set) + mismatch badge
 - link to open call review/evaluation

Final outcome

- editable by:
 - agent? (optional)
 - manager? (only if admin allows)
 - admin always
- shows “final outcome confidence” if AI predicts mismatch

Attribution details

- UTM/campaign payload (if imported)
- “Manual attribution” fallback

Contact-based AI behavior (v1)

- call type guess improves:
 - “This is the 3rd call with this number → likely follow-up”
 - outcome mismatch detection improves:
 - “Contact already marked lost → why did agent mark converted?”
-

Part 6 — Call Ingestion journeys (deep detail)

A) Upload ingestion (Bulk)

Actors: Admin/Manager (primary), Agent (optional)

Step-by-step

1. User opens Upload
2. Drops audio files
3. Uploads CSV mapping
4. System validates:
 - filenames match
 - agent exists or can be created/mapped
 - datetime parse
 - language field valid
5. System shows preview table:
 - row count
 - warning count
 - errors count
6. User fixes errors or proceeds
7. Processing begins:
 - each call shows status:
 - queued → transcribing → scoring → ready
8. Completion:
 - success list

- failed list with retry button
- “unassigned agent” list with assign action

Engineering must support

- idempotency (re-upload doesn't duplicate)
 - retry on failure
 - partial success
-

B) Exotel ingestion

Actors: Admin sets up, system auto-ingests, manager reviews mapping

Step-by-step

1. Admin connects Exotel
2. System starts ingesting recordings
3. Calls arrive into Calls Inbox:
 - may be “Unassigned agent”
4. Manager/admin maps agent identity:
 - mapping table: extension/number/name → agent user
5. Once mapped, future calls auto-assign agent

UX required: Agent mapping table

- shows “unknown agent identifiers” detected
 - dropdown to assign to an agent user
 - save mapping
 - optionally apply retroactively to historical calls
-

C) Twilio ingestion

Same as Exotel, but with Twilio-specific fields behind the scenes.

Part 7 — Notifications and nudges (v1)

In-app notifications

- New evaluation available (agent)
- Pending acknowledgement reminder (agent)
- Clarification request received (manager)
- Outcome mismatch flagged (manager)
- Integration error (admin)
- Weekly insights ready (manager/admin)

Email notifications (v1)

- Invite email
- Weekly insights email (optional toggle)
- Pending acknowledgements digest (optional toggle)

Design requirement: Notification preferences page.

Part 8 — Exports and Evidence Packs (v1)

Evidence Pack (shareable link)

- Generated from call review page
- Includes:
 - score summary
 - per-criterion results
 - evidence snippets with timestamps
 - transcript excerpt
 - call metadata
- Shareable link:
 - expiry date (user-chosen)
 - “copy link” action
 - revoke link action

Export formats supported (your requirement: all)

- PDF evidence pack
- CSV evaluations
- audio download
- CRM-note copy block (structured summary)

Edge cases

- transcript redaction (optional v1 toggle)
 - missing audio (should block evidence pack)
-

Part 9 — What is “done” in v1 (boundaries)

To avoid scope creep, v1 is complete when:

Must-have

- Upload + Exotel + Twilio ingestion
- Transcript + audio sync + speaker labeling
- Scorecard builder + templates + versioning
- Auto scoring with evidence + confidence + unclear
- QA queue workflow
- Agent evaluations + acknowledgement + clarification
- Call type AI guess + manual override
- Outcomes per scorecard + final outcome per contact journey
- Outcome mismatch flagging
- Insights:
 - ROI trends by source/channel
 - QA trends
 - Voice of customer basics
- Evidence pack export + share link expiry
- Roles/permissions + admin toggles for manager editing
- Hindi + English + Punjabi support
- SLA: scoring results visible within 5 minutes (for typical call lengths)

Explicitly later

- Missed calls
 - Full AI coaching learning plans
 - Mobile app
 - Multi-location workspaces
-

Part 10 — Design Team To-Do (v1)

This is the checklist your designers can literally work through.

A) System-wide UX artifacts

- Information architecture per role (nav differences)
- Component library:
 - tables with sticky headers
 - filter chips
 - “status pills” (processing, unclear, fail, pass)
 - evidence snippet cards
 - transcript search UI
- Empty states (critical):
 - no calls yet
 - no scorecard yet
 - no agents yet
 - no matched attribution yet
- Loading states:
 - ingestion queue
 - transcript generation
 - scoring pending (<5 min)
- Error states:
 - transcription failed
 - missing audio
 - CSV validation errors
 - integration disconnected
- Notification UX:
 - in-app toast + notification center
 - email templates

B) New screens not yet fully “designed”

(We discussed them conceptually; these need explicit UI)

- Contact Journey page (timeline + final outcome editing)
- Attribution settings page (source/channel mapping + lead metadata import)
- Agent mapping table for Exotel/Twilio unknown agent identifiers
- Outcome mapping UI (call outcome → journey stage)

C) Micro-interactions

- “Jump to evidence” behavior (play audio + highlight transcript)
 - “Override score” confirmation pattern
 - “Request clarification” modal flow
 - “Outcome mismatch” warning prompt
-

Part 11 — Engineering Team To-Do (v1)

This is the build plan disguised as a checklist.

A) Foundations

- Auth (invite-based)
- Roles (Admin/Manager/Agent)
- Team membership
- Permission toggles (admin-configurable)

B) Data model (minimum)

- Workspace, Team, User
- Call:
 - audio asset
 - metadata
 - transcript + diarization
 - call type (ai + manual)
 - scorecard id + version
 - score results (per criterion)
 - flags
 - agent-set call outcome + mismatch flag
- Contact:
 - phone number
 - calls list
 - final outcome + stage
 - attribution payload
- Scorecard:
 - sections
 - criteria
 - outcomes per scorecard
 - versions
- Mapping:
 - call outcome → journey stage/final outcome

C) Ingestion pipelines

- Upload + CSV ingestion:
 - CSV parsing + validation
 - file matching
 - idempotency
- Exotel ingestion:
 - fetch recordings + metadata
 - agent mapping fallback
- Twilio ingestion:
 - fetch recordings + metadata
 - agent mapping fallback

D) Processing pipeline (must hit <5 mins)

- async job queue architecture
- transcription (multi-language)
- diarization / speaker labeling
- scoring engine:
 - criterion-level outputs
 - evidence snippet extraction
 - confidence scoring
 - unclear state
- indexing for:
 - transcript search
 - keyword triggers
 - VoC clustering

E) Workflows

- QA Queue generation rules:
 - unclear
 - low score
 - critical fail
 - new rep
 - random sample
 - trigger terms
- Override storage + recalculation
- Agent acknowledgement workflow
- Clarification workflow

F) Reporting and insights

- Source/channel aggregation
- QA trends computation
- VoC extraction:
 - objections
 - lost reasons
 - competitor mentions
- Contact journey funnel reporting

G) Exports

- PDF evidence pack generator
- Shareable link system with expiry + revoke
- CSV export
- Audio download access control
- CRM note generator

H) Observability

- ingestion success rate
- transcription success rate
- scoring SLA monitoring
- unclear rate monitoring
- override rate monitoring (model drift signal)